

## 13.4 Deletion in red-black trees

Deletion in a red-black tree is similar to insertion.


- Apply the deletion algorithm for binary search trees.
- Apply node color changes and left/right rotations to fix the violations of RBT tree properties.

Suppose we delete a node  $z$ , and there are three cases of  $z$  in a binary search tree:

(i)  $z$  is a leaf; (ii)  $z$  has only one child; and (iii)  $z$  has two children.

When  $z$  has two children, let  $y$  be the successor of  $z$  in the RBT, the actual deletion proceeds as follows: Copy the content of  $y$  (not its color) to node  $z$ , node  $y$  then is “spliced out” (check slides 17-18 in Lecture 12 about the deletion operation in a binary search tree), where  $y$  only has only the right child in the binary search tree.

## 13.1 Deletion in Binary Search Trees



Deletion in binary search trees (Cormen et al., p297)

## 13.4 Deletion in red-black trees

We now adjust the colors of some nodes in the tree.

- If  $y$  was RED, the new tree is already a red-black tree, done.
- If  $y$  was BLACK, the new tree due to the removal of  $y$  has several types of red-black property violations:
  - Case one: If  $y$  was the tree root, the new root might be RED
  - Case two: If both the parent and the child of  $y$  are RED, they now have the relationship between the parent and the child
  - Case three: Any path passing through  $y$  now has fewer BLACK nodes.

## 13.4 Deletion in red-black trees

Suppose the color of  $y$  was BLACK and  $x$  is the sole child of  $y$  in the tree (Why?).

- If  $y$  was the root and a red child  $x$  of  $y$  now becomes the new root, then Property 2 of the RBT is violated.
- If both  $x$  and  $p[y]$  are RED before the deletion of  $y$ , then Property 4 of the RBT is violated.
- The removal of  $y$  causes any path containing  $y$  has a black node less. This causes Property 5 of the RBT to be violated.

We correct the above errors by saying that node  $x$  has an “extra” BLACK color.

## 13.4 Deletion in red-black trees

To restore Property 5 of the RBT, we consider four cases:

- Case 1:  $x$ 's sibling  $w$  is red
- Case 2:  $x$ 's sibling  $w$  is black, and both children of  $w$  are black
- Case 3:  $x$ 's sibling  $w$  is black,  $w$ 's left child is red and  $w$ 's right child is black
- Case 4:  $x$ 's sibling  $w$  is black, and the right child of  $w$  is red

**The key idea is that in each case the number of black nodes from the subtree root to each of the subtrees  $\alpha, \beta, \dots$  is preserved by the transformation.**

The general strategy: (a) Transform Case 1 to Case 2 and Case 3 to Case 4. (b) Solve Case 2 and Case 4, respectively.

## 13.4 Deletion in red-black trees

Case 1:  $x$ 's sibling  $w$  is red.

Since  $w$  must have black children (Why?),

- swap the colors of  $w$  and  $p(x)$ ,
- perform a left-rotation on  $p(x)$  (without violating any of red-black tree properties)

**The new sibling of  $x$ , which is one of the  $w$ 's children prior to the rotation, now is black.** Thus, Case 1 has been converted to one of Cases 2, 3, and 4.

Cases 2, 3, and 4 occur when  $w$  is black, they are distinguished by the colors of  $w$ 's children.

## 13.4 Deletion in red-black trees

Case 2:  $x$ 's sibling  $w$  is black, and both children of node  $w$  are black.

Since  $w$  is black, take one black off from both  $x$  and  $w$  (as both children of  $w$  are black), leaving  $x$  with only one black and leaving (or coloring)  $w$  red.

To compensate for removing one black from  $x$  and  $w$ , we add an extra black to  $p(x)$ , which originally either red or black. If  $p(x)$  was red, now it is black, done; otherwise, node  $p(x)$  now becomes new node  $x$ , i.e.,  $p(x)$  now has double black on it.

Notice that node  $x$  (**with double black**) has been pushed one-layer higher if  $p(x)$  was black already, where a layer near to the tree root is a higher layer. Otherwise, done.

The above procedure continues until (i) it reaches the tree root; (ii) or is converted into one of the other three cases. If the tree root becomes double black, remove the extra black, and Property 5 of the RBT is still held.

**This case takes  $O(h)$  time where  $h$  is the height of the RBT.**

## 13.4 Deletion in red-black trees

Case 3:  $x$ 's sibling  $w$  is black,  $w$ 's left child is red while  $w$ 's right child is black.

- swap the colors of  $w$  and its left child  $left(w)$ ,
- perform a right rotation on  $w$  without violating any of the red-black tree properties.

The verification of the red-black tree properties can be done by checking whether any path passing through node  $w$  still contains the same number of black nodes from its source node to any of its descendant leaves before and after this transformation.

The new sibling  $w$  of  $x$  now is black with a red right child, this becomes Case 4.



## 13.4 Deletion in red-black trees


Case 4:  $x$ 's sibling  $w$  is black, and the right child of  $w$  is red.

Remove the extra black represented by  $x$ , by changing some colors and performing a left rotation on node  $p(x)$ , the extra black of  $x$  can be removed, making it a single black without violating any red-black tree properties.

- swap the colors of  $w$  and its right child  $right(w)$ ,
- swap colors of  $p(x)$  and  $w$
- perform the left rotation on  $p(x)$ ,
- color  $x$ 's parent  $p(x)$  with black (i.e., remove the extra black from  $x$ ).

The above operations do not violating any of the red-black tree properties.

## 13.4 Deletion in red-black trees



Four cases in fixing red-black errors during deletion (Cormen et al., p329)

## 13.4 Analysis on Deletion in red-black trees

**Theorem:** Given a RBT containing  $N$  nodes, the removal of one node from it takes  $O(\log N)$  time if the resulting tree is still a RBT.

**The sketch of the proof:** Since the height of any red-black tree is  $O(\log N)$ , each of Cases 1, 3 and 4 takes  $O(1)$  time by performing a left or right rotation and re-coloring, while Case 2 takes  $O(h)$  time, where  $h$  is the height of the red-black tree and  $h = O(\log N)$  by the lemma in slides 5-6 of Lecture 13. Thus, the deletion operation takes  $O(\log N)$  time.